

```
In [1]:
```

```
import keras
keras.__version__
```

```
Using TensorFlow backend.
```

```
Out[1]:
```

```
'2.2.4'
```

5.2 - Using convnets with small datasets

This notebook contains the code sample found in Chapter 5, Section 2 of [Deep Learning with Python](#). Note that the original text features far more content, in particular further explanations and figures: in this notebook, you will only find source code and related comments.

Training a convnet from scratch on a small dataset

Having to train an image classification model using only very little data is a common situation, which you likely encounter yourself in practice if you ever do computer vision in a professional context.

Having "few" samples can mean anywhere from a few hundreds to a few tens of thousands of images. As a practical example, we will focus on classifying images as "dogs" or "cats", in a dataset containing 4000 pictures of cats and dogs (2000 cats, 2000 dogs). We will use 2000 pictures for training, 1000 for validation, and finally 1000 for testing.

In this section, we will review one basic strategy to tackle this problem: training a new model from scratch on what little data we have. We will start by naively training a small convnet on our 2000 training samples, without any regularization, to set a baseline for what can be achieved. This will get us to a classification accuracy of 71%. At that point, our main issue will be overfitting. Then we will introduce *data augmentation*, a powerful technique for mitigating overfitting in computer vision. By leveraging data augmentation, we will improve our network to reach an accuracy of 82%.

In the next section, we will review two more essential techniques for applying deep learning to small datasets: *doing feature extraction with a pre-trained network* (this will get us to an accuracy of 90% to 93%), and *fine-tuning a pre-trained network* (this will get us to our final accuracy of 95%). Together, these three strategies -- training a small model from scratch, doing feature extracting using a pre-trained model, and fine-tuning a pre-trained model -- will constitute your future toolbox for tackling the problem of doing computer vision with small datasets.

The relevance of deep learning for small-data problems

You will sometimes hear that deep learning only works when lots of data is available. This is in part a valid point: one fundamental characteristic of deep learning is that it is able to find interesting features in the training data on its own, without any need for manual feature engineering, and this can only be achieved when lots of training examples are available. This is especially true for problems where the input samples are very high-dimensional, like images.

However, what constitutes "lots" of samples is relative -- relative to the size and depth of the network you are trying to train, for starters. It isn't possible to train a convnet to solve a complex problem with just a few tens of samples, but a few hundreds can potentially suffice if the model is small and well-regularized and if the task is simple. Because convnets learn local, translation-invariant features, they are very data-efficient on perceptual problems. Training a convnet from scratch on a very small image dataset will still yield reasonable results despite a relative lack of data, without the need for any custom feature engineering. You will see this in action in this section.

But what's more, deep learning models are by nature highly repurposable: you can take, say, an image classification or speech-to-text model trained on a large-scale dataset then reuse it on a significantly different problem with only minor changes. Specifically, in the case of computer vision, many pre-trained models (usually trained on the ImageNet dataset) are now publicly available for download and can be used to bootstrap powerful vision models out of very little data. That's what we will do in the next section.

For now, let's get started by getting our hands on the data.

Downloading the data

The cats vs. dogs dataset that we will use isn't packaged with Keras. It was made available by Kaggle.com as part of a computer

vision competition in late 2013, back when convnets weren't quite mainstream. You can download the original dataset at: <https://www.kaggle.com/c/dogs-vs-cats/data> (you will need to create a Kaggle account if you don't already have one -- don't worry, the process is painless).

The pictures are medium-resolution color JPEGs. They look like this:



Unsurprisingly, the cats vs. dogs Kaggle competition in 2013 was won by entrants who used convnets. The best entries could achieve up to 95% accuracy. In our own example, we will get fairly close to this accuracy (in the next section), even though we will be training our models on less than 10% of the data that was available to the competitors. This original dataset contains 25,000 images of dogs and cats (12,500 from each class) and is 543MB large (compressed). After downloading and uncompressing it, we will create a new dataset containing three subsets: a training set with 1000 samples of each class, a validation set with 500 samples of each class, and finally a test set with 500 samples of each class.

Here are a few lines of code to do this:

In [2]:

```
import os, shutil
```

In [7]:

```
# The path to the directory where the original
# dataset was uncompressed
# update dataset directory
original_dataset_dir = 'E:/code_nlp/assign03/cats_and_dogs/train'
original_dataset_dir = '/Users/ly/Desktop/train'

# The directory where we will
# store our smaller dataset
# base_dir = 'E:/code_nlp/assign03/small'
# update dataset directory
base_dir = '/Users/ly/Desktop/small'
os.mkdir(base_dir)

# Directories for our training,
# validation and test splits
train_dir = os.path.join(base_dir, 'train')
os.mkdir(train_dir)

validation_dir = os.path.join(base_dir, 'validation')
os.mkdir(validation_dir)
test_dir = os.path.join(base_dir, 'test')
os.mkdir(test_dir)

# Directory with our training cat pictures
train_cats_dir = os.path.join(train_dir, 'cats')
```

```

train_cats_dir = os.path.join(train_dir, 'cats')
os.mkdir(train_cats_dir)

# Directory with our training dog pictures
train_dogs_dir = os.path.join(train_dir, 'dogs')
os.mkdir(train_dogs_dir)

# Directory with our validation cat pictures
validation_cats_dir = os.path.join(validation_dir, 'cats')
os.mkdir(validation_cats_dir)

# Directory with our validation dog pictures
validation_dogs_dir = os.path.join(validation_dir, 'dogs')
os.mkdir(validation_dogs_dir)

# Directory with our validation cat pictures
test_cats_dir = os.path.join(test_dir, 'cats')
os.mkdir(test_cats_dir)

# Directory with our validation dog pictures
test_dogs_dir = os.path.join(test_dir, 'dogs')
os.mkdir(test_dogs_dir)

# Copy first 1000 cat images to train_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(train_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copy next 500 cat images to validation_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copy next 500 cat images to test_cats_dir
fnames = ['cat.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_cats_dir, fname)
    shutil.copyfile(src, dst)

# Copy first 1000 dog images to train_dogs_dir
fnames = ['dog.{}.jpg'.format(i) for i in range(1000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(train_dogs_dir, fname)
    shutil.copyfile(src, dst)

# Copy next 500 dog images to validation_dogs_dir
fnames = ['dog.{}.jpg'.format(i) for i in range(1000, 1500)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(validation_dogs_dir, fname)
    shutil.copyfile(src, dst)

# Copy next 500 dog images to test_dogs_dir
fnames = ['dog.{}.jpg'.format(i) for i in range(1500, 2000)]
for fname in fnames:
    src = os.path.join(original_dataset_dir, fname)
    dst = os.path.join(test_dogs_dir, fname)
    shutil.copyfile(src, dst)

```

As a sanity check, let's count how many pictures we have in each training split (train/validation/test):

In [8]:

```
print('total training cat images:', len(os.listdir(train_cats_dir)))
```

total training cat images: 1000

In [9]:

```
print('total training dog images:', len(os.listdir(train_dogs_dir)))
```

total training dog images: 1000

In [10]:

```
print('total validation cat images:', len(os.listdir(validation_cats_dir)))
```

total validation cat images: 500

In [11]:

```
print('total validation dog images:', len(os.listdir(validation_dogs_dir)))
```

total validation dog images: 500

In [12]:

```
print('total test cat images:', len(os.listdir(test_cats_dir)))
```

total test cat images: 500

In [13]:

```
print('total test dog images:', len(os.listdir(test_dogs_dir)))
```

total test dog images: 500

So we have indeed 2000 training images, and then 1000 validation images and 1000 test images. In each split, there is the same number of samples from each class: this is a balanced binary classification problem, which means that classification accuracy will be an appropriate measure of success.

Building our network

We've already built a small convnet for MNIST in the previous example, so you should be familiar with them. We will reuse the same general structure: our convnet will be a stack of alternated `Conv2D` (with `relu` activation) and `MaxPooling2D` layers.

However, since we are dealing with bigger images and a more complex problem, we will make our network accordingly larger: it will have one more `Conv2D + MaxPooling2D` stage. This serves both to augment the capacity of the network, and to further reduce the size of the feature maps, so that they aren't overly large when we reach the `Flatten` layer. Here, since we start from inputs of size 150x150 (a somewhat arbitrary choice), we end up with feature maps of size 7x7 right before the `Flatten` layer.

Note that the depth of the feature maps is progressively increasing in the network (from 32 to 128), while the size of the feature maps is decreasing (from 148x148 to 7x7). This is a pattern that you will see in almost all convnets.

Since we are attacking a binary classification problem, we are ending the network with a single unit (a `Dense` layer of size 1) and a `sigmoid` activation. This unit will encode the probability that the network is looking at one class or the other.

In [14]:

```
from keras import layers
from keras import models

model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
```

```
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))
```

WARNING: Logging before flag parsing goes to stderr.

W0716 18:51:41.819855 4509238720 deprecation_wrapper.py:119] From /anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:74: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

W0716 18:51:41.905030 4509238720 deprecation_wrapper.py:119] From /anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

W0716 18:51:41.935769 4509238720 deprecation_wrapper.py:119] From /anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:4138: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

W0716 18:51:41.991240 4509238720 deprecation_wrapper.py:119] From /anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3976: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_pool2d instead.

Let's take a look at how the dimensions of the feature maps change with every successive layer:

In [15]:

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 148, 148, 32)	896
max_pooling2d_1 (MaxPooling2D)	(None, 74, 74, 32)	0
conv2d_2 (Conv2D)	(None, 72, 72, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 36, 36, 64)	0
conv2d_3 (Conv2D)	(None, 34, 34, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 17, 17, 128)	0
conv2d_4 (Conv2D)	(None, 15, 15, 128)	147584
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 128)	0
flatten_1 (Flatten)	(None, 6272)	0
dense_1 (Dense)	(None, 512)	3211776
dense_2 (Dense)	(None, 1)	513
Total params: 3,453,121		
Trainable params: 3,453,121		
Non-trainable params: 0		

Explanation for deriving unknown parameters

The number 896 in layer conv2d_1 is calculated as follows: since the inputs are colorful images (depth of 3), the number of dendrites is $k \times k = 9$ from every neuron in every one of 32 layers of conv2d_1, due to parameter sharing, we have $3 \times 9 \times 32$ weights plus one bias for every filter layer $3 \times 9 \times 32 + 32 = 896$.

The number 73856 in layer conv2d_3 is calculated as:

1. Every one of 128 filter layers in conv2d_3 "convolves" with every one of 64 output layers of maxpooling2d_2. There are: $128 \times 64 = 8192$ layer to layer pairs.
2. All neurons within a filter layer (one of 128 layers) in conv2d_3 are the same and have $k \times k = 9$ unknown weights (CNNs parameter sharing) on dendrites pointing to every one of 64 input layers. Thus we have $8192 \times 9 = 73728$ unknown weights

weights.

3. Each one of 128 layers in conv2d_3 has its own bias parameter, and we have $73728 + 128 = 73856$ unknown parameters.

For our compilation step, we'll go with the `RMSprop` optimizer as usual. Since we ended our network with a single sigmoid unit, we will use binary crossentropy as our loss (as a reminder, check out the table in Chapter 4, section 5 for a cheatsheet on what loss function to use in various situations).

In [16]:

```
from keras import optimizers
```

```
model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

```
W0716 18:51:45.892283 4509238720 deprecation_wrapper.py:119] From
/anaconda3/envs/pytf/lib/python3.6/site-packages/keras/optimizers.py:790: The name
tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
W0716 18:51:45.900501 4509238720 deprecation_wrapper.py:119] From
/anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3376: The nam
e tf.log is deprecated. Please use tf.math.log instead.
```

```
W0716 18:51:45.909523 4509238720 deprecation.py:323] From /anaconda3/envs/pytf/lib/python3.6/site-
packages/tensorflow/python/ops/nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from
tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
```

Instructions for updating:

Use `tf.where` in 2.0, which has the same broadcast rule as `np.where`

Data preprocessing

As you already know by now, data should be formatted into appropriately pre-processed floating point tensors before being fed into our network. Currently, our data sits on a drive as JPEG files, so the steps for getting it into our network are roughly:

- Read the picture files.
- Decode the JPEG content to RGB grids of pixels.
- Convert these into floating point tensors.
- Rescale the pixel values (between 0 and 255) to the [0, 1] interval (as you know, neural networks prefer to deal with small input values).

It may seem a bit daunting, but thankfully Keras has utilities to take care of these steps automatically. Keras has a module with image processing helper tools, located at `keras.preprocessing.image`. In particular, it contains the class `ImageDataGenerator` which allows to quickly set up Python generators that can automatically turn image files on disk into batches of pre-processed tensors. This is what we will use here.

In [17]:

```
from keras.preprocessing.image import ImageDataGenerator
```

```
# All images will be rescaled by 1./255
```

```
train_datagen = ImageDataGenerator(rescale=1./255)
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

```
train_generator = train_datagen.flow_from_directory(
```

```
    # This is the target directory
```

```
    train_dir,
```

```
    # All images will be resized to 150x150
```

```
    target_size=(150, 150),
```

```
    batch_size=20,
```

```
    # Since we use binary_crossentropy loss, we need binary labels
```

```
    class_mode='binary')
```

```
validation_generator = test_datagen.flow_from_directory(
```

```
    validation_dir,
```

```
    target_size=(150, 150),
```

```
    batch_size=20,
```

```
    class_mode='binary')
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

Let's take a look at the output of one of these generators: it yields batches of 150x150 RGB images (shape (20, 150, 150, 3)) and binary labels (shape (20,)). 20 is the number of samples in each batch (the batch size). Note that the generator yields these batches indefinitely: it just loops endlessly over the images present in the target folder. For this reason, we need to `break` the iteration loop at some point.

We need to install pillow package. Pillow is a PIL fork. PIL is Python Image Library

In [19]:

```
!pip install pillow
```

Requirement already satisfied: pillow in /anaconda3/lib/python3.7/site-packages (5.3.0)

In [20]:

```
for data_batch, labels_batch in train_generator:
    print('data batch shape:', data_batch.shape)
    print('labels batch shape:', labels_batch.shape)
    break
```

```
data batch shape: (20, 150, 150, 3)
labels batch shape: (20,)
```

Let's fit our model to the data using the generator. We do it using the `fit_generator` method, the equivalent of `fit` for data generators like ours. It expects as first argument a Python generator that will yield batches of inputs and targets indefinitely, like ours does. Because the data is being generated endlessly, the generator needs to know example how many samples to draw from the generator before declaring an epoch over. This is the role of the `steps_per_epoch` argument: after having drawn `steps_per_epoch` batches from the generator, i.e. after having run for `steps_per_epoch` gradient descent steps, the fitting process will go to the next epoch. In our case, batches are 20-sample large, so it will take 100 batches until we see our target of 2000 samples.

When using `fit_generator`, one may pass a `validation_data` argument, much like with the `fit` method. Importantly, this argument is allowed to be a data generator itself, but it could be a tuple of Numpy arrays as well. If you pass a generator as `validation_data`, then this generator is expected to yield batches of validation data endlessly, and thus you should also specify the `validation_steps` argument, which tells the process how many batches to draw from the validation generator for evaluation.

In [21]:

```
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=15,
    validation_data=validation_generator,
    validation_steps=50)
```

```
W0716 18:54:35.712226 4509238720 deprecation_wrapper.py:119] From
/anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:986: The name
tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.
```

```
Epoch 1/15
100/100 [=====] - 32s 321ms/step - loss: 0.6864 - acc: 0.5405 - val_loss:
0.6632 - val_acc: 0.6020
Epoch 2/15
100/100 [=====] - 33s 333ms/step - loss: 0.6595 - acc: 0.6045 - val_loss:
0.6451 - val_acc: 0.6170
Epoch 3/15
100/100 [=====] - 32s 317ms/step - loss: 0.6089 - acc: 0.6580 - val_loss:
0.6206 - val_acc: 0.6530
Epoch 4/15
100/100 [=====] - 33s 331ms/step - loss: 0.5612 - acc: 0.7025 - val_loss:
0.6369 - val_acc: 0.6370
Epoch 5/15
100/100 [=====] - 32s 325ms/step - loss: 0.5403 - acc: 0.7285 - val_loss:
0.6301 - val_acc: 0.6750
Epoch 6/15
100/100 [=====] - 32s 324ms/step - loss: 0.5210 - acc: 0.7400 - val_loss:
0.6201 - val_acc: 0.6850
```



```

100/100 [=====] - 33s 334ms/step - loss: 0.5210 - acc: 0.7420 - val_loss:
0.6425 - val_acc: 0.6630
Epoch 7/15
100/100 [=====] - 33s 334ms/step - loss: 0.5040 - acc: 0.7495 - val_loss:
0.5941 - val_acc: 0.6820
Epoch 8/15
100/100 [=====] - 32s 321ms/step - loss: 0.4829 - acc: 0.7600 - val_loss:
0.5665 - val_acc: 0.7050
Epoch 9/15
100/100 [=====] - 32s 323ms/step - loss: 0.4565 - acc: 0.7780 - val_loss:
0.5690 - val_acc: 0.7000
Epoch 10/15
100/100 [=====] - 32s 321ms/step - loss: 0.4276 - acc: 0.7985 - val_loss:
0.6227 - val_acc: 0.6960
Epoch 11/15
100/100 [=====] - 33s 332ms/step - loss: 0.4059 - acc: 0.8145 - val_loss:
0.5854 - val_acc: 0.7030
Epoch 12/15
100/100 [=====] - 32s 324ms/step - loss: 0.3846 - acc: 0.8285 - val_loss:
0.5444 - val_acc: 0.7260
Epoch 13/15
100/100 [=====] - 32s 315ms/step - loss: 0.3570 - acc: 0.8495 - val_loss:
0.5879 - val_acc: 0.7120
Epoch 14/15
100/100 [=====] - 32s 322ms/step - loss: 0.3401 - acc: 0.8550 - val_loss:
0.6011 - val_acc: 0.7150
Epoch 15/15
100/100 [=====] - 32s 324ms/step - loss: 0.3121 - acc: 0.8765 - val_loss:
0.5591 - val_acc: 0.7360

```

It is good practice to always save your models after training:

In [22]:

```
!pip install h5py
```

```

Requirement already satisfied: h5py in /anaconda3/lib/python3.7/site-packages (2.8.0)
Requirement already satisfied: numpy>=1.7 in /anaconda3/lib/python3.7/site-packages (from h5py)
(1.15.4)
Requirement already satisfied: six in /anaconda3/lib/python3.7/site-packages (from h5py) (1.11.0)

```

In [23]:

```
model.save('cats_and_dogs_small_1.h5')
```

Let's plot the loss and accuracy of the model over the training and validation data during training:

In [26]:

```

import matplotlib.pyplot as plt

acc = history.history['acc']
val_acc = history.history['val_acc']
loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))


plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

```

In [28]:

```
plt.show()
```

These plots are characteristic of overfitting. Our training accuracy increases linearly over time, until it reaches nearly 100%, while our validation accuracy stalls at 70-72%. Our validation loss reaches its minimum after only five epochs then stalls, while the training loss keeps decreasing linearly until it reaches nearly 0.

Because we only have relatively few training samples (2000), overfitting is going to be our number one concern. You already know about a number of techniques that can help mitigate overfitting, such as dropout and weight decay (L2 regularization). We are now going to introduce a new one, specific to computer vision, and used almost universally when processing images with deep learning models: *data augmentation*.

Using data augmentation

Overfitting is caused by having too few samples to learn from, rendering us unable to train a model able to generalize to new data. Given infinite data, our model would be exposed to every possible aspect of the data distribution at hand: we would never overfit. Data augmentation takes the approach of generating more training data from existing training samples, by "augmenting" the samples via a number of random transformations that yield believable-looking images. The goal is that at training time, our model would never see the exact same picture twice. This helps the model get exposed to more aspects of the data and generalize better.

In Keras, this can be done by configuring a number of random transformations to be performed on the images read by our `ImageDataGenerator` instance. Let's get started with an example:

In [29]:

```
datagen = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest')
```

These are just a few of the options available (for more, see the Keras documentation). Let's quickly go over what we just wrote:

- `rotation_range` is a value in degrees (0-180) a range within which to randomly rotate pictures

- `rotation_range` is a value in degrees (0-180), a range within which to randomly rotate pictures.
- `width_shift` and `height_shift` are ranges (as a fraction of total width or height) within which to randomly translate pictures vertically or horizontally.
- `shear_range` is for randomly applying shearing transformations.
- `zoom_range` is for randomly zooming inside pictures.
- `horizontal_flip` is for randomly flipping half of the images horizontally -- relevant when there are no assumptions of horizontal asymmetry (e.g. real-world pictures).
- `fill_mode` is the strategy used for filling in newly created pixels, which can appear after a rotation or a width/height shift.

Let's take a look at our augmented images:

In [30]:

```
# This is module with image preprocessing utilities
from keras.preprocessing import image

fnames = [os.path.join(train_cats_dir, fname) for fname in os.listdir(train_cats_dir)]

# We pick one image to "augment"
img_path = fnames[7]



# Read the image and resize it
img = image.load_img(img_path, target_size=(150, 150))



# Convert it to a Numpy array with shape (150, 150, 3)
x = image.img_to_array(img)

# Reshape it to (1, 150, 150, 3)
x = x.reshape((1,) + x.shape)

# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in datagen.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 4 == 0:
        break

plt.show()
```





Pick another picture as an example for data augmentation

Define data generator and Pick file 17

In [49]:

```
rotation = ImageDataGenerator(
    rotation_range=40,
    width_shift_range=0,
    height_shift_range=0,
    shear_range=0,
    zoom_range=0,
    horizontal_flip=False,
    fill_mode='nearest')

widthChange = ImageDataGenerator(
    rotation_range=0,
    width_shift_range=0.5,
    height_shift_range=0,
    shear_range=0,
    zoom_range=0,
    horizontal_flip=False,
    fill_mode='nearest')

htShift = ImageDataGenerator(
    rotation_range=0,
    width_shift_range=0,
    height_shift_range=0.3,
    shear_range=0,
    zoom_range=0,
    horizontal_flip=False,
    fill_mode='nearest')

sheer = ImageDataGenerator(
    rotation_range=0,
    width_shift_range=0,
    height_shift_range=0,
    shear_range=0.3,
    zoom_range=0,
    horizontal_flip=False,
    fill_mode='nearest')

zoom = ImageDataGenerator(
    rotation_range=0,
    width shift range=0,
```

```

height_shift_range=0,
shear_range=0,
zoom_range=0.3,
horizontal_flip=False,
fill_mode='nearest')

hFlip = ImageDataGenerator(
    rotation_range=0,
    width_shift_range=0,
    height_shift_range=0,
    shear_range=0,
    zoom_range=0,
    horizontal_flip=True,
    fill_mode='nearest')

# We pick one image to "augment"
img_path = fnames[17]

# Read the image and resize it
img = image.load_img(img_path, target_size=(150, 150))

# Convert it to a Numpy array with shape (150, 150, 3)
x = image.img_to_array(img)

# Reshape it to (1, 150, 150, 3)
x = x.reshape((1,) + x.shape)

```

Apply rotation

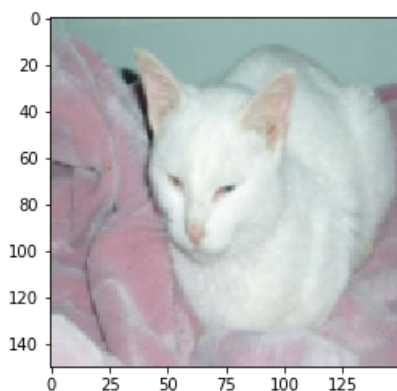

In [50]:

```

# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in rotation.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 2 == 0:
        break

plt.show()



```



Apply width change

In [51]:

```
# The .flow() command below generates batches of randomly transformed images.  
# It will loop indefinitely, so we need to `break` the loop at some point!  
i = 0  
for batch in widthChange.flow(x, batch_size=1):  
    plt.figure(i)  
    imgplot = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 2 == 0:  
        break  
  
plt.show()
```




Apply height shift

In [52]:

```
# The .flow() command below generates batches of randomly transformed images.  
# It will loop indefinitely, so we need to `break` the loop at some point!  
i = 0  
for batch in htShift.flow(x, batch_size=1):  
    plt.figure(i)  
    imgplot = plt.imshow(image.array_to_img(batch[0]))  
    i += 1  
    if i % 2 == 0:  
        break  
  
plt.show()
```



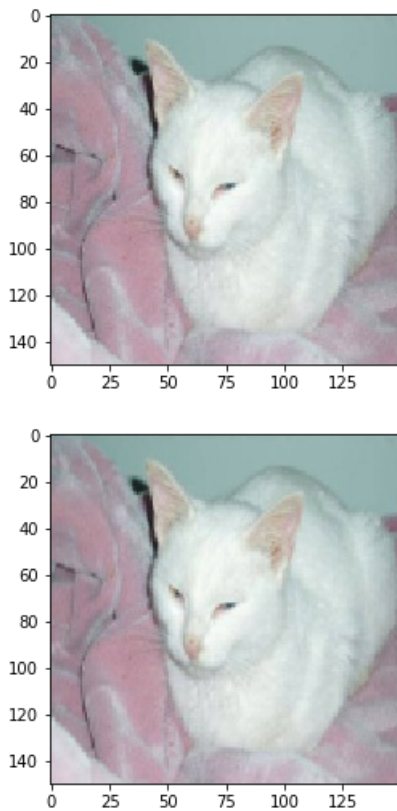


Apply shear

In [53]:

```
# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in shear.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 2 == 0:
        break

plt.show()
```



Apply zoom

In [54]:



```
# The .flow() command below generates batches of randomly transformed images
```

```

# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in zoom.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 2 == 0:
        break

plt.show()

```



Apply horizon flip


In [55]:


```

# The .flow() command below generates batches of randomly transformed images.
# It will loop indefinitely, so we need to `break` the loop at some point!
i = 0
for batch in hFlip.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i += 1
    if i % 2 == 0:
        break

plt.show()

```





If we train a new network using this data augmentation configuration, our network will never see twice the same input. However, the inputs that it sees are still heavily intercorrelated, since they come from a small number of original images -- we cannot produce new information, we can only remix existing information. As such, this might not be quite enough to completely get rid of overfitting. To further fight overfitting, we will also add a Dropout layer to our model, right before the densely-connected classifier:

In [31]:

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu',
                        input_shape=(150, 150, 3)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Flatten())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy',
              optimizer=optimizers.RMSprop(lr=1e-4),
              metrics=['acc'])
```

W0716 19:11:18.944018 4509238720 deprecation.py:506] From /anaconda3/envs/pytf/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Let's train our network using data augmentation and dropout:

In [32]:

```
train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=40,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,)

# Note that the validation data should not be augmented!
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    # This is the target directory
    train_dir,
    # All images will be resized to 150x150
    target_size=(150, 150),
    batch_size=32,
    # Since we use binary_crossentropy loss, we need binary labels
    class_mode='binary')
```

```
validation_generator = test_datagen.flow_from_directory(
    validation_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='binary')
```

```
history = model.fit_generator(
    train_generator,
    steps_per_epoch=100,
    epochs=15,
    validation_data=validation_generator,
    validation_steps=50)
```

Found 2000 images belonging to 2 classes.

Found 1000 images belonging to 2 classes.

Epoch 1/15

100/100 [=====] - 53s 529ms/step - loss: 0.6916 - acc: 0.5188 - val_loss: 0.6801 - val_acc: 0.5114

Epoch 2/15

100/100 [=====] - 54s 537ms/step - loss: 0.6776 - acc: 0.5669 - val_loss: 0.6622 - val_acc: 0.6108

Epoch 3/15

100/100 [=====] - 49s 494ms/step - loss: 0.6549 - acc: 0.6037 - val_loss: 0.6558 - val_acc: 0.6079

Epoch 4/15

100/100 [=====] - 49s 493ms/step - loss: 0.6387 - acc: 0.6297 - val_loss: 0.6086 - val_acc: 0.6508

Epoch 5/15

100/100 [=====] - 51s 510ms/step - loss: 0.6241 - acc: 0.6444 - val_loss: 0.6022 - val_acc: 0.6732

Epoch 6/15

100/100 [=====] - 51s 508ms/step - loss: 0.6101 - acc: 0.6697 - val_loss: 0.5767 - val_acc: 0.6920

Epoch 7/15

100/100 [=====] - 50s 496ms/step - loss: 0.6010 - acc: 0.6663 - val_loss: 0.5596 - val_acc: 0.7157

Epoch 8/15

100/100 [=====] - 53s 529ms/step - loss: 0.5863 - acc: 0.6825 - val_loss: 0.5661 - val_acc: 0.6875

Epoch 9/15

100/100 [=====] - 51s 510ms/step - loss: 0.5837 - acc: 0.6922 - val_loss: 0.5482 - val_acc: 0.7101

Epoch 10/15

100/100 [=====] - 52s 523ms/step - loss: 0.5681 - acc: 0.7028 - val_loss: 0.5523 - val_acc: 0.7069

Epoch 11/15

100/100 [=====] - 53s 527ms/step - loss: 0.5643 - acc: 0.7106 - val_loss: 0.5281 - val_acc: 0.7300

Epoch 12/15

100/100 [=====] - 51s 508ms/step - loss: 0.5717 - acc: 0.6969 - val_loss: 0.5407 - val_acc: 0.7107

Epoch 13/15

100/100 [=====] - 58s 578ms/step - loss: 0.5567 - acc: 0.7084 - val_loss: 0.5327 - val_acc: 0.7262

Epoch 14/15

100/100 [=====] - 56s 559ms/step - loss: 0.5491 - acc: 0.7119 - val_loss: 0.5689 - val_acc: 0.6980

Epoch 15/15

100/100 [=====] - 53s 527ms/step - loss: 0.5467 - acc: 0.7197 - val_loss: 0.5477 - val_acc: 0.7113

Let's save our model -- we will be using it in the section on convnet visualization.

In [34]:

```
model.save('cats_and_dogs_small_2.h5')
```

Let's plot our results again:

In [35]:

```
acc = history.history['acc']
val_acc = history.history['val_acc']
```

```

loss = history.history['loss']
val_loss = history.history['val_loss']

epochs = range(len(acc))



plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.legend()

plt.figure()

plt.plot(epochs, loss, 'bo', label='Training loss')
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.legend()

plt.show()

```



Thanks to data augmentation and dropout, we are no longer overfitting: the training curves are rather closely tracking the validation curves. We are now able to reach an accuracy of 82%, a 15% relative improvement over the non-regularized model.

By leveraging regularization techniques even further and by tuning the network's parameters (such as the number of filters per convolution layer, or the number of layers in the network), we may be able to get an even better accuracy, likely up to 86-87%. However, it would prove very difficult to go any higher just by training our own convnet from scratch, simply because we have so little data to work with. As a next step to improve our accuracy on this problem, we will have to leverage a pre-trained model, which will be the focus of the next two sections.